

## TP4 - WEB-based Distributed Applications

Organisation	Binome
Evaluation	Code-based
Advised time	6-8h per binome
Tools	JDK ==1.8, Eclipse
Return	<b>Eclipse Projet</b> , named <b>LastName1.LastName2</b> . In compressed format, either zip or .tar ( <b>nothing else</b> )
Contact	<a href="mailto:Fabienne.Boyer@univ-grenoble-alpes.fr">Fabienne.Boyer@univ-grenoble-alpes.fr</a> <a href="mailto:Olivier.Gruber@univ-grenoble-alpes.fr">Olivier.Gruber@univ-grenoble-alpes.fr</a>

### 1. Objective

This practical work aims at programming a Web Server over TCP/IP. You will have to design and implement a Web Server able to serve static HTML pages as well as dynamic requests that involve *ricmlets*. Ricmlets are closed to regular HTTP servlets, just simpler. Ricmlets provide clients with the ability to run Java code, server side, following a well-defined API.

To be usable from any browser, your Web Server will follow the specifications of the HTTP protocol. However **for time reasons, we will only consider serving the GET requests of the HTTP protocol**. As you all know, GET requests allow to get either static pages or dynamic ones, as described below.

### 2. HTTP Protocol Basics

Let's consider the basics of the HTTP protocol. The protocol recognizes HTTP requests and HTTP responses, both using the same generic message format ([RFC7230](https://tools.ietf.org/html/rfc7230)). We will consider a simplified version described below:

```
HTTP-message grammar = start-line
                        *( header-field CRLF )
                        CRLF
                        [ message-body ]
```

Let's discuss the different parts:

- A start-line present in all HTTP message
- Zero or more header fields, each header field being terminated by the two characters CR and LF
- An empty line marking the end of the header
- Optionally, a message body as a sequence of bytes

Overall, the message contains a superset of printable US-ASCII characters. Special characters are CR LF and SP. The character **CR** is the US-ASCII value 13 (0xD), the character **LF** is the US-ASCII value 10 (0xA). The character **SP** (space) is the ASCII value 32 (0x20).

The start-line has either a request-line for requests or a status-line for responses.

*start-line* = *request-line* | *status-line*

*request-line* = *method* SP *request-target* SP *HTTP-version* CRLF (ex: GET /hello.html HTTP/1.1)

*status-line* = HTTP-version SP *status-code* SP *reason-phrase* CRLF (ex: HTTP/1.1 200 OK)

Header fields are lines beginning with a field name, followed by a colon (":"), followed by a field body, and terminated by CRLF.

header-field = field-name : field-body. (ex: Content-type: text/html)

### 3. Serving static Pages

Client-Side Protocol:

- The client connects to the server, via an TCP/IP socket
- The client sends a request with a start-line indicating a method and a resource identifier. The resource identifier may be followed by the version of the HTTP protocol used by the client. For instance, to get a static page named `hello.html` server side, the start-line may be:

```
GET /hello.html HTTP/1.1
```

Server-Side Protocol:

- The server accepts on a given port number for connections from remote clients.
- For any accepted connection, the server creates a dedicated worker thread (for simplicity, we advise using a basic multi-threaded design, but feel free to decide otherwise).
- When receiving a request, the server looks for a file corresponding to the requested page.
- If the file is found, the server returns a response composed of a start-line indicating a success, followed by a header indicating the length and the type of the body. After the header is sent, the server sends the file content.

```
HTTP/1.1 200 OK
Date: ...
Server: Ricm4HttpServer
Content-length: 45876
Content-type: text/html
```

- To build the response header, you may get the content type through the suffix of the resource identifier given in the url (html, txt, gif, jpg, etc). Just use the static method `getContentType(String)` that is given to you in the class *HttpRequest*.
- In case the requested file is not found at server side, the server must reply with an error code

```
HTTP/1.0 404 File not found
```

Work to do:

- **Understand** the interfaces *HttpRequest* and *HttpResponse* in the package *httpserver.itf*.
- Then **understand** the class *HTTPServer* in the Java package *httpserver.impl* and how it uses the implementation classes *HttpRequest* and *HttpResponseImpl*. Together, these classes implement a web server who serves HTTP GET requests and serves the corresponding static HTML pages. To get started, look at the entry point (static method *main*) on the class *HTTPServer*.
- Then **complete the class *HttpStaticRequest*** that handles static HTTP requests.
- Check your implementation with a local browser first: try to get the url <localhost:<port>/hello.html>, assuming that the file *hello.html* is accessible from the file folder of your server. Try to get different kinds of resources, *gif*, *jpeg*, or others (we give you a FILES folder containing examples of such files within the Eclipse project).
- Check also your HTTP server from a terminal, through executing the *wget* command (install *wget* if required).
- If you have time, check your HTTP server from a remote browser through the url <hostname:<port>/hello.html>.

### 4. Dynamic Pages

We now consider dynamic pages, meaning that clients requests can trigger the execution of server-side classes (that we call *ricmlets*) that return dynamically computed HTML pages. As for Java servlets, a ricmlet should define a method `doGet(..)` that takes as argument request and response objects. Overall, the protocol at client and server sides is the following.

**Protocol at client side**

- The client connects to the server.
- The client sends a request with a start-line indicating a method and a resource identifier. The resource starts by */ricmlets* and indicates the name of the ricmlet to execute. For instance, to execute a ricmlet implemented by the class *examples.HelloRicmlet*, the start-line will be:

```
GET /ricmlets/examples/HelloRicmlet HTTP/1.1
```

- In case your ricmlet expects arguments, these are sent along with the resource identifier, as it is the case with regular servlets:

```
GET /ricmlets/examples/HelloRicmlet?name=Bob&surname=Marley HTTP/1.1
```

**Protocol at server side**

- The server waits on a given port number for connections from remote clients.
- When receiving a request, the HTTP server determines if it is a static or dynamic request. In case of dynamic request, **it looks for the class corresponding to the requested ricmlet under the *ricmlets* package.**
- If the class is found, the server instantiates it if it is not already instantiated. It is very important to respect this lifecycle management constraint: in this current step, we want that each ricmlet behaves as a singleton, meaning that there should exist at most one instance per ricmlet class.
- Once the ricmlet instance has been found, the server executes the *doGet()* method, giving as arguments two objects: one representing the request and another one representing the response. Among others, the *request* object allows to get the arguments, the *response* object allows to get the stream on which the response body should be sent. See the class *HelloRicmlet.java* for an example of a ricmlet class.
- In case the Ricmlet class is not found at server side, the server replies with an error code

```
HTTP/1.0 404 Ricmlet not found
```

Notice that ricmlet classes do not set up the *content-length* header field in their response. They produce HTML responses dynamically and incrementally, the bytes composing the response being directly written on the server output stream for efficiency reasons. It is thus difficult to predetermine the length of the response before producing the response itself. Hopefully, not indicating the length of the response is not an issue for HTML responses, because browsers have the ability to detect when the full response has been received through HTML tags.

**Work to do**

- Complete your *HttpServer* class to support dynamic requests (in a very first step, you may not consider arguments). We ask you to define two additional classes: *HttpRicmletRequestImpl* and *HttpRicmletResponseImpl*, that respectively extends *HttpRequest* and *HttpResponseImpl*.

Remind that you can instantiate a class whose name is stored in a string as follow (assuming the class has a no-argument public constructor):

```
String s = "a.b.c.Foo";
Class.forName(s).newInstance();
```

- Check your implementation from a browser: try to get the url <localhost:<port>/ricmlets/examples/HelloRicmlet>.
- Check that your ricmlet lifecycle management is correct. To this end, use the ricmlet *CountRicmlet* that delivers a page indicating the number of times it has been processed. Check that getting the url <localhost:<port>/ricmlets/examples/CountRicmlet> several times actually increments this number.
- Then consider managing arguments, through implementing the *getArgs()* method in the class *HttpRicmletRequest*. Check with your browser that the url <localhost:<port>/ricmlets/examples/HelloRicmlet?name=Bob&surname=Marley> delivers a page including a *Hello Bob Marley* message.

**5. Session Management**

We now consider *sessions*, a concept discussed during the lecture on HTTP. **We will firstly consider non-persistent session, meaning sessions that do not survive a kill-restart cycle on your *HttpServer*.**

You will integrate the notion of session through relying on *cookies* managed by the HTTP protocol. The HTTP protocol includes a special *Set-Cookie* command requesting browsers to keep a cookie at client side. A response header may include several *Set-Cookie* commands. Moreover, a *Set-cookie* command may define several cookies, as illustrated below.

```
HTTP/1.0 200 OK
...
Set-Cookie: myFirstCookie=123, mySecondCookie=Hello
Set-Cookie: anotherCookie=45678
Server: Ricm4HttpServer
Content-length: 45876
Content-type: text/html
```

Reversely, once a browser keeps a cookie sent by a server *S*, the browser will automatically integrate this cookie in any request sent to the server *S*, in the header part of the request, as illustrated below:

```
GET /ricmlets/example/HelloRicmlet HTTP/1.1
...
Cookie: myFirstCookie=123
Cookie: mySecondCookie=Hello
Cookie: anotherCookie=45678
...
```

Coming back to session management, the idea will be represent a session by an object. The *HttpServer* will be in charge of creating such objects and associate unique identifiers to them. It will also be in charge of integrating the identifier of a session, as a cookie, in any response sent to a browser.

### Work to do

- Complete your *HttpServer* classes to support sessions. In a first step, do not manage of session destructions, just focus on session creations. Define a class *Session* that implements the given *HttpSession* interface. Then instantiate this class at the right moment in your Server implementation.
- Check your implementation through the ricmlet *CountBySessionRicmlet* that counts the number of times it has been processed per session. To this end, fetch the following url from two distinct client sessions, checking that each session gets its proper counter:  
<http://localhost:<port>/ricmlets/examples/CountBySessionRicmlet>
- Pay attention to act as two users such as to get two distinct sessions. Try with opening two browsers (e.g., *firefox* and *chrome*) on your machine. If your configuration forces these two browsers to share a same cookie space, you may use a local browser and a remote one. You may also use a local browser and a local terminal running the *wget* command.
- **[OPTIONAL]** Then manage session destructions. Remember that a session gets automatically destroyed in case the client does not issue any request for a given duration.

## 6. Application Management

We finally extend our Web Server with the notion of *applications* to give a way to modularize ricmlet classes. We will consider that an application defines a set of ricmlet classes through a Java archive file. All archive files will be put at a given folder at server-side. Then, urls used to trigger the execution of ricmlets will follow the pattern given below:

[<hostname>:<port>/ricmlets/<appName>/<fullRicmletClassName>](http://<hostname>:<port>/ricmlets/<appName>/<fullRicmletClassName>).

Upon such request, the server will look for the class *<ricmletClassName>* in the *<appName>.jar* file. Then it will get the classloader associated to the application *appName* and request it to instantiate this class if not yet done. As we have seen together during the lecture, using a classloader per application is the only way to ensure that an application cannot instantiate a class belonging to another application.

### Work to do

- Define a class *Application* that provides a method *getInstance(String className, String appName, ClassLoader parent)* allowing to get the instance of the ricmlet associated to the class *className*, for the account of the application *appName*.
- Adapt your *HttpServer* classes to take into account applications.

- Check your implementation. For this, make a jar file named *app1.jar*, containing the ricmlet *CountRicmlet* for instance. Then, fetch the url `http://localhost:<port>/app1/CountRicmlet` from your browser and check that the count printed by *CountRicmlet* increases each time you fetch the url.
- Then make another jar file named *app2.jar*, containing the same ricmlets classes. Then, fetch the url `http://localhost:<port>/app2/CountRicmlet` from your browser. You should observe that the count printed by *CountRicmlet* starts at 0.